

Paper Number(s): **E2.19**

IMPERIAL COLLEGE LONDON

DEPARTMENT OF ELECTRICAL AND ELECTRONIC ENGINEERING
EXAMINATIONS 2009

EEE Part II: MEng, BEng and ACGI

INTRODUCTION TO COMPUTER ARCHITECTURE

Monday, 1 June 2.00 pm

Time allowed: 1:30 hours

Corrected Copy



There are FOUR questions on this paper.

Question 1 is compulsory and carries 40% of the marks.

Answer Question 1 and two others from Questions 2-4 which carry equal marks (30% each).

Any special instructions for invigilators and information for candidates are on page 1.

Examiners responsible:

First Marker(s): Clarke, T.
Second Marker(s): Demir, Y.

Special information for invigilators:

The booklet Exam Notes 2009 should be distributed with the Examination Paper.

Information for candidates:

The prefix &, or suffix $_{(16)}$, introduces a hexadecimal number, e.g: &1C0, 1C0 $_{(16)}$.

The booklet Exam Notes 2009, as published on the course web pages, is provided and contains reference material.

Question 1 is compulsory and carries 40% of marks. Answer only TWO of the Questions 2-4, which carry equal marks.

The Questions

1. [Compulsory]

- (a) Perform the following numeric conversions:
- (i) $-251_{(10)}$ into 9 bit signed octal
 - (ii) "abcA" into 32 bit hexadecimal ASCII codes with the first character stored in the least significant byte.
 - (iii) $-22.25_{(10)}$ into IEEE-754 floating point. Give your answer in *hexadecimal*.
- [8]
- (b) A CPU with 60 MHz clock frequency executes one instruction per cycle using a pipeline of length of 5, and has a pipeline stall every 10 cycles. Calculate the average throughput of the CPU in millions of instructions per second (MIPS), assuming stall time = pipeline length. What would be the throughput of this CPU with no pipeline?
- [8]
- (c) State the value in *decimal* of registers R0-R5 at the end of the ARM assembly code fragment in *Figure 1.1* assuming that at the start of the code fragment $Rn = n$ ($n = 0, 1, \dots, 14$).
- [12]
- (d)
- (i) State the unsigned and signed number ranges for an n bit binary number.
 - (ii) State in the ARM architecture what is the Boolean expression on condition codes **N,Z,C,V** for unsigned arithmetic overflow.
 - (iii) State in the ARM architecture what is the Boolean expression on condition codes **N,Z,C,V** that will implement the condition $R0 > R1$ after instruction **CMP R0,R1** assuming R0, R1 are signed.
- [12]

```
RSBS    R0, R7, R6
ADD     R1, R6, R7, lsr #1
EOR     R2, R7, R8
SBC     R3, R3, R3
BIC     R4, R10, #3
MOV     R5, R12, lsl #10
```

Figure 1.1

2. Each code fragment (a) - (c) below executes with all condition codes and registers initially 0, and memory locations as in *Figure 2.1*. State the value of R0-R3, the condition codes, and any *changed* memory locations, after execution of the code fragment. Write your answers using as a template a copy of the table in *Figure 2.3* omitting the row labelled (x) which indicates the required format of your answer. Each answer must be written in hexadecimal, except the condition codes which must be in binary, this format illustrated in row (x).

- (a) Code as in *Figure 2.2a*. [10]
- (b) Code as in *Figure 2.2b*. [10]
- (c) Code as in *Figure 2.2c*. [10]

Location	Value
&100	&11121314
&104	&10203040
&108	&01020304
&10C	&80706050
> &10C	&0

Figure 2.1 - memory locations

```
MOV    R10, #&100
MOV    R11, #4
LDR    R0, [R10,R11]
LDR    R1, [R10,#8]!
LDRB   R2, [R10],#1
LDRB   R3, [R10]
STRB   R10, [R11,R11]
```

(a)

```
MOV    R0, #&108
MOV    R1, #&200
LDMDA  R0!, {R2,R3}
STMIB  R1, {R2,R3}
```

(b)

```
MOV    R0, #1
MOV    R1, R0, rol #10
EORS   R2, R1, R0, ror #1
ADC    R3, R0, R0
SUBS   R0, R0, R0
```

(c)

Figure 2.2 - code fragments

	R0	R1	R2	R3	NZCV	Memory
(x)	0	&1020	&FFFFFFF	&C	0110	mem ₈ [&120] = &10 mem ₃₂ [&300] = &FFFF0000
(a)						
(b)						
(c)						

Figure 2.3 - template for answers

3. The ARM code in *Figure 3.1* sets R1 to a value which depends on R2, R3, R4.

- (a) Give code examples from the execution of this code to show how an ARM instruction that enters the FETCH stage of the pipeline may be either not executed, condition-true executed, or condition-false executed.

[6]

- (b) If initially $R1 = x1$, $R2 = x2$, $R3 = x3$, $R4 = x4$, state concisely, using pseudo-code with one or more if-then-else statements, what is the final value to which R1 is set.

[8]

- (c) If R2,R3,R4 are initially 0 trace through the execution of the ARM7 code in *Figure 3.1* illustrating in a diagram the instructions occupying each stage of the pipeline in every cycle.

[8]

- (d) State what are the quickest and slowest paths through the code in *Figure 3.1* when executed on the ARM7, giving in each case the code execution time in machine cycles. Instruction timing may be found in the Exam Notes booklet.

[8]

```
A    CMP R2, #0
B    ADDGE R1, R1, R2
C    SUBLT R1, R1, R2
D    BEQ X
E    ADD R1, R1, R3
F    B Y
X    ADD R1, R1, R4
Y
```

Figure 3.1

4. This question relates to the ARM assembler code fragment LOOP in *Figure 4.1*.

- (a) If R0 has non-negative value n at the start of LOOP, calculate as a function of n the number of iterations of LOOP. Discuss what happens if n is negative.

[6]

- (b) The instructions with labels C & D test ARM condition codes. State in each case which instruction sets the condition codes which are tested, and therefore what is the condition on data in R1 for each of these instructions to be condition-true executed.

[8]

- (c) Suppose that just before the instruction with label A is executed, the bits of R1 & R3 are denoted X(31:0) and Y(31:0). Using these bit designations, determine the value of each bit of R3 just after the instruction with label D is executed.

[8]

- (d) At the start of LOOP, $R5 = p$, $R6 = q$. Determine the addresses of all memory locations loaded and stored in the i th iteration of LOOP, where i ranges from 1 upwards, as a function of i , p and q . Hence state concisely what is the change in memory locations made by this code when it is executed with $R0 = n$.

[8]

```
LOOP
    LDR    R1, [R5], #4
    LDR    R3, [R6]
A    BIC    R3, R3, #&80000003
B    ANDS   R4, R1, #&80000001
C    EORMI  R3, R3, R4
D    ORRNE  R3, R3, #2
    STR    R3, [R6], #4
    SUBS   R0, R0, #1
    BGE    LOOP
```

Figure 4.1

EXAM NOTES 2009

Introduction to Computer Architecture (EE2)

Introduction to Computer Architecture and Systems (ISE)

Memory Reference & Transfer Instructions

LDR load word
STR store word
LDRB load byte
STRB store byte
LDREQB ; note position
: of EQ
STREQB

LDMEID r13!, {r0-r4,r6,r7} ; ! => write-back to register
STMFA r13, {r2}
STMEQB r2!, {r5-r12} ; note position of EQ
; higher reg nos go to/from higher mem addresses always
[EIP][AID] empty/full, ascending/descending
[IID][AIB] inc/decr, after/before

LDR r0, [r1]
LDR r0, [r1, #offset]
LDR r0, [r1, #offset]!
LDR r0, [r1], #offset
LDR r0, [r1, r2]
LDR r0, [r1, r2, #shift]
LDR r0, address_label
ADR r0, address_label
; register-indirect addressing
; pre-indexed addressing
; pre-indexed, auto-indexing
; post-indexed, auto-indexing
; register-indexed addressing
; scaled register-indexed addressing
; PC relative addressing
; load PC relative address

R2.1

Conditions Binary Encoding

Opcode [31:28]	Mnemonic extension	Interpretation	Status flag state for execution
0000	EQ	Equal / equals zero	Z set
0001	NE	Not equal	Z clear
0010	CS/HS	Carry set / unsigned higher or same	C set
0011	CC/LO	Carry clear / unsigned lower	C clear
0100	MI	Minus / negative	N set
0101	PL	Plus / positive or zero	N clear
0110	VS	Overflow	V set
0111	VC	No overflow	V clear
1000	HI	Unsigned higher	C set and Z clear
1001	LS	Unsigned lower or same	C clear or Z set
1010	GE	Signed greater than or equal	N equals V
1011	LT	Signed less than	N is not equal to V
1100	GT	Signed greater than	Z clear and N equals V
1101	LE	Signed less than or equal	Z set or N is not equal to V
1110	AL	Always	any
1111	NV	Never (do not use!)	none

R2.2

ARM Data Processing Instructions Binary Encoding

Opcode [24:21]	Mnemonic	Meaning	Effect
0000	AND	Logical bit-wise AND	Rd := Rn AND Op2
0001	EOR	Logical bit-wise exclusive OR	Rd := Rn EOR Op2
0010	SUB	Subtract	Rd := Rn - Op2
0011	RSB	Reverse subtract	Rd := Op2 - Rn
0100	ADD	Add	Rd := Rn + Op2
0101	ADC	Add with carry	Rd := Rn + Op2 + C
0110	SBC	Subtract with carry	Rd := Rn - Op2 + C - 1
0111	RSC	Reverse subtract with carry	Rd := Op2 - Rn + C - 1
1000	TST	Test	See on Rn AND Op2
1001	TEQ	Test equivalence	See on Rn EOR Op2
1010	CMP	Compare	See on Rn - Op2
1011	CMN	Compare negated	See on Rn + Op2
1100	ORR	Logical bit-wise OR	Rd := Rn OR Op2
1101	MOV	Move	Rd := Op2
1110	BIC	Bit clear	Rd := Rn AND NOT Op2
1111	MVN	Move negated	Rd := NOT Op2

R2.3

Data Processing Operand 2

Examples

ADD r0, r1, r2
MOV r0, #1
CMP r0, #-1
EOR r0, r1, r2, lsr #10
RSB r0, r1, r2, asr r3

Op2	Conditions	Notes
Rm		
#imm	imm = s rotate 2r (0 ≤ s ≤ 255, 0 ≤ r ≤ 15)	Assembler will translate negative values changing op-code as necessary Assembler will work out rotate if it exists
Rm, shift #s	(1 ≤ s ≤ 31)	rx always sets carry
Rm, rrx #1	shift => lsr, lsl, asr, ror	ror sets carry if S=1
Rm, shift Rs	shift => lsr, lsl, asr, ror	shifts do not set carry
		shift by register value (takes 2 cycles)

R2.4

Multiply Instructions

- MUL, MLA were the original (32 bit result) instructions
 - Why does it not matter whether they are signed or unsigned?
- Later architectures added 64 bit results
- Note that some multiply instructions have 4 register operands!
 - Multiply instructions must have register operands, no immediate constant
 - Multiplication by small constants can often be implemented more efficiently with data processing instructions – see Lecture 10.

NB d & m must be different for MUL, MLA

ARM3 and above

MUL rd, rm, rs multiply (32 bit) $Rd := (Rm * Rs)[31:0]$
 MLA rd, rm, rs, rn multiply-acc (32 bit) $Rd := (Rm * Rs)[31:0] + Rn$
 UMLAL r1, r1, rm, rs unsigned multiply $(R1, R1) := R1 * Rs$
 UMLAL r1, r1, rm, rs unsigned multiply-acc $(R1, R1) := (R1 * Rs) + R1$
 SMLAL r1, r1, rm, rs signed multiply $(R1, R1) := R1 * Rs$
 SMLAL r1, r1, rm, rs signed multiply-acc $(R1, R1) := (R1 * Rs) + R1$
ARM7DM core and above

Exceptions & Interrupts

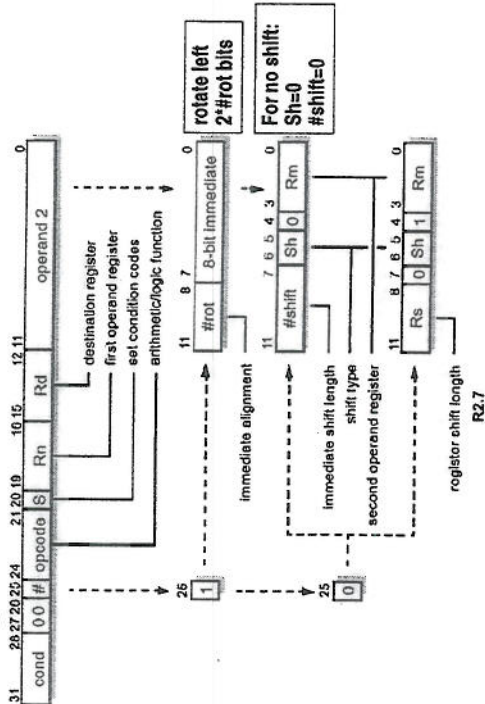
Exception	Return
SWI or undefined instruction	MOVSPC, R14
IRQ, FIQ, prefetch abort	SUBSPC, R14, #4
Data abort (needs to return failed instruction)	SUBSPC, R14, #8

Exception Mode	Shadow registers
SVC, UND, IRQ, Abort	R13, R14, SPSR
FIQ	as above + R8-R12

(0x introduces a hex constant)

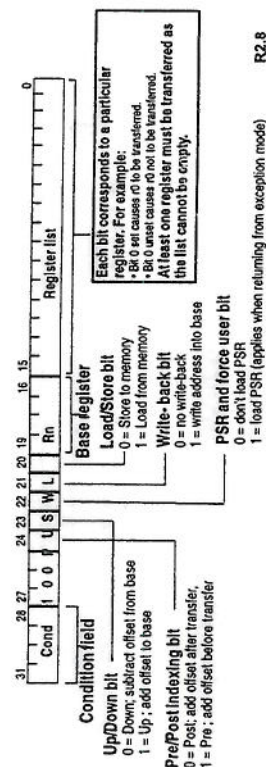
Exception	Mode	Vector address
Reset	SVC	0x00000000
Undefined instruction	UND	0x00000004
Software interrupt (SWI)	SVC	0x00000008
Prefetch abort (instruction fetch memory fault)	Abort	0x0000000C
Data abort (data access memory fault)	Abort	0x00000010
IRQ (normal interrupt)	IRQ	0x00000018
FIQ (fast interrupt)	FIQ	0x0000001C

Data Processing Instruction Binary Encoding




Multiple Register Transfer Binary Encoding

- The Load and Store Multiple instructions (LDM / STM) allow between 1 and 16 registers to be transferred to or from memory.



Branch Instruction Binary Encoding


❖ Branch: $B\{\langle \text{cond} \rangle\} \text{ label}$ ❖ Branch with Link: `BL{<cond>}` `sub_routine_label`

- ❖ The offset for branch instructions is calculated by the assembler:

- + By taking the difference between the branch instruction and the target address minus 8 (to allow for the pipeline).
- + This gives a 26 bit offset which is right shifted 2 bits (as the bottom two bits are always zero as instructions are word – aligned) and stored into the instruction encoding.
- + This gives a range of ± 32 Mbytes.

R2.9

Instruction Set Overview



ARM Instruction Timing

Exact Instruction timing is very complex and depends in general on memory cycle times which are system dependent. The table below gives an approximate guide.

Instruction	Typical execution time (cycles) (If instruction condition is TRUE – otherwise 1 cycle)
Any instruction, with condition false	1
data processing (all except when shift is by number equal to value of register)	1
data processing (register-valued shifts)	2
LDR, LDRB	4
STR, STRB	4
LDM (n registers)	n+3 (+3 if PC is loaded)
STM (n registers)	n+3
B, BL	4
Multiply	7-14 (varies with architecture & operand values)

ASCII Codes

NB - b7 = 0

b(3:0)

0	NUL	SOH	STX	ETX	EDT	EHQ	ACK	BEL	BS	HT	LF	VT	FF	CR	SO	SI
1	DLE	DC1	DC2	DC3	DC4	NAK	SYN	ETB	CAN	EM	SUB	ESC	F9	GB	RS	US
2	SPC	!	"	#	\$	%	&	'	()	*	+	,	-	.	/
3	0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
4	@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5	P	Q	R	S	T	U	V	W	X	Y	Z	[\]	^	_
6	`	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
7	p	q	r	s	t	u	v	w	x	y	z	{		}	~	DEL

A=analysis, D=design, C=calculated solution using taught theory, B=bookwork

NB - marking will be 1 mark for every 2% indicated on question paper (max 50 marks) so marks here are half of marks on question paper.

Solution to Question 1

36 minutes for the question => 9 minutes each part

(a)

i) $405_{(8)}$

ii) $61626341_{(16)}$ (NB assume little-endian so LSB is bits (7:0) of the word)

iii) $C1B20000_{(16)}$ $s=1$, $e=131$, $m=(1.)011001$

[1 mark each, except iii 2 marks]

[4C]

(b)

$T=60M/(1+0.1*5)=40MIPS$. With no pipeline there is no stall, but stages happen sequentially so $60M/5 = 12MIPS$

[4C]

(c)

$R0 = -1$,

$R1 = 6+7/2 = 9$,

$R2 = 15$,

$R3 = -1$ (no carry from RSBS),

$R4 = 8$,

$R5 = 12*1024 = 12288$

[6C]

d)

(i) $0-2^n-1$ (unsigned), $-2^{n-1} - 2^{n-1}-1$ (signed)

(ii) C

(iii) either result is positive & no overflow, or result is negative & overflow:

$!(N).!V + N.V).!(Z)$ or

$!(N \oplus V).!Z$ etc

[6A]

Solution to Question 2

27 minutes for the question

This tests ability to understand low-level operation of ARM assembler instructions

For each part, deduct 1 mark for each column wrong down to minimum of 0 marks. Assume $\text{mem}[] = \text{mem}_{32}[]$.

	r0	r1	r2	r3	NZCV	Memory
a)	&10203040	&01020304	&04	&03	n/a	$\text{mem}_8[\&8] = \&108$
b)	&100	&200	&10203040	&01020304	n/a	$\text{mem}_{32}[\&204] = \&10203040$ $\text{mem}_{32}[\&208] = \&01020304$
c)	&0	&400	&80000400	&2	0110	n/a

Note ;

consequent errors allowed (e.g. data wrong in memory write)

In b) r2/r3 swapped => 1 mark

[5A+5A+5A]

Solution to Question 3

27 minutes for the question

This questions tests understanding of instruction execution in the ARM architecture.

(a)

Not executed but FETCHED: instruction X if F is executed

condition-true executed: A

condition-false executed: B if $R2 < 0$ (signed)

[3B]

b)

if $x2 > 0$ (signed) then $R1 := x1 + x2 + x3$ elseif $x2 < 0$ (signed) $R1 := x1 - x2 + x3$ if $x2 = 0$ then $R1 := x1 + x2 + x4$

[4A]

c)

FETCH	A	B	C	D	E	F	wait	X	Y	...
DECODE		A	B	C	D	E	stall	stall	X	Y
EXECUTE			A	B	C	D	stall	stall	stall	X

[4A/B]

d)

quickest: A,B,C,D,X (8 cycles)

slowest: A,B,C,D,E,F (9 cycles)

[4A]

Solution to Question 4

This question tests whether the student understands the ARM bit manipulation & conditional execution.

27 minutes for the question

a)

$n + 1$ iterations (loop for values of n down to and including 0). If n is negative it will loop once.

[3A]

b) Instruction at B sets codes for both C & D. C & V are not set, N & Z are set based on the value of R3 AND $\&80000001$ (bitwise AND).

Hence C is executed if $R1(31)$ is 1 \Leftrightarrow R1 negative, D if $R1(31)$ or $R1(0)$ is 1.

[4A]

c) There are three bits of R3 which can change: 31, 1 and 0.

Instructions A, C & D may change these as follows

	31	1	0	
A	0	0	0	
C	$R1(31)$	0	$R1(31).R1(0)$	change if $R1(31)=true$
D	$R1(31)$	$R1(31)+R1(0)$	$R1(31).R1(0)$	change if $R1(31)$ or $R1(0)=true$

Hence

$$R3(31) = R1(31)$$

$$R3(1) = R1(31)+R1(0)$$

$$R3(0) = R1(31).R1(0)$$

All other bits of R3 stay the same.

(NB $R1=X$, $R3=Y$)

[4A]

d)

In the first iteration the addresses are:

R1 load: p

R3 load: q

R3 store: q

Each iteration these both advance by 4, so we have

R1 load $p+4(i-1)$

R3 load and store $q+4(i-1)$

this continues for $n+1$ iterations $i= 1$ to $n+1$

Hence $n+1$ words $[q, q+4, \dots, q+4n]$ have bits 31,1,0 modified as per part (c) by the corresponding bits of corresponding words in $[p, p+4, \dots, p+4n]$. If the two memory areas overlap with $q > p$ the changes to q affect subsequent p locations.

[4A]