

Question 1(a) *New computed example*

[R= Running, - = waiting]

Round Robin with time slice of 3 ms

A	R	R	R																	
B			-	R	R															
C						R	R	R												
D							-	-	R	R	R	-	-	R	R					
E							-	-	-	-	-	R	R							
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19

AWT= (0+1+0+4+4)/5= 1.8ms; ATT= 24/5 = 4.8ms

Advantages: Simple to implement and fair

Disadvantages: Difficult to determine appropriate time quantum:

- Too small: good response time, but large overheads (scheduler is called too often)
- Too large: bad response time

[3] [2 for correct solution, + 1 for adv/disadvantages]**After marking comments: most people got the execution of the algorithm correctly, with answers to the advantages/disadvantages variable, but mostly ok**

Priority scheduling with pre-emption

A	R	R	-	-	R															
B			R	R																
C						R	-	-	-	-	-	-	-	R	R					
D							R	-	-	R	R	R	R							
E								R	R											
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19

AWT=11/5= 2.2ms ATT= 26/5 = 5.2ms

Advantages:

Takes into account external factors regarding the importance of the various processes

Disadvantages:

Might result (quite probable) in starvation of low-priority processes – this can be avoided using an *aging* procedure; processes that wait for too long have their priorities gradually increased.**[2] [2; 1 for correct solution, + 1 for adv/disadvantages]****After marking comments: most people got the execution of this algorithm correctly too, with answers to the advantages/disadvantages variable, but mostly ok.**

SJF

A	R	R	R																	
B			-	R	R															
C						R	R	R												
D							-	-	-	-	R	R	R	R	R					
E								-	R	R										
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19

Avg waiting time: (0+1+0+4+1) / 5 = 1.2 ms, Avg turnaround time: 21/5 = 4.2 ms

[2; 1 for correct solution, + 1 for adv/disadvantages]**Advantages:** Provably optimal in that it gives the minimum average waiting time for a given set of processes**Disadvantages:** Knowing the length of next CPU burst is difficult; frequently this is predicted utilising previous lengths as estimates, or is user-specified; can result in long waits (including starvation) for long processes.**After marking comments: most people got the execution of this algorithm correctly too, with answers to the advantages/disadvantages variable, but mostly ok. Some people did not note that the algorithm is provably optimal.**

(b) [new computed example]

Optimal page replacement algorithm (7 page faults)

	1	4	5	6	9	4	1	5	4	7
Frame1	1	1	1	1	1			5		7
Frame2	-	4	4	4	4			4		4
Frame3	-	-	5	6	9			9		9

There are multiple correct ways of doing the last two replacements [2]

FIFO page replacement algorithm (9 page faults)

	1	4	5	6	9	4	1	5	4	7
Frame1	1	1	1	6	6	6	1	1		1
Frame2	-	4	4	4	9	9	9	5		5
Frame3	-	-	5	5	5	4	4	4		7

[2]

LRU (Least recently used) page replacement algorithm (9 page faults)

	1	4	5	6	9	4	1	5	4	7
Frame1	1	1	1	6	6	6	1	1		7
Frame2	-	4	4	4	9	9	9	5		5
Frame3	-	-	5	5	5	4	4	4		4

[2]

After marking comments: most people got the execution of these algorithms correctly, particularly the FIFO. Some minor mistakes with Optimal and LRU, with no particular pattern in the errors.

(c) [Bookwork]

Turn is a character variable; Interested_A and Interested_B are boolean variables initially set to FALSE;

```
Interested_A = TRUE;
Turn = 'B';
while (interested_B = TRUE
      AND Turn = 'B')
    Do_nothing;
Critical_A;
Interested_A = FALSE;
```

```
Interested_B = TRUE;
Turn = 'A';
while (interested_A = TRUE
      AND Turn = 'A')
    Do_nothing;
Critical_B;
interested_B = FALSE;
```

[1 point for setting the interest variables, 1 point for setting the turn variables, 1 point for setting the loop conditions correctly and 1 point for setting the interest variables to false after [4]

After marking comments: most of those that attempted this got it correctly; some errors included forgetting to give the turn to the OTHER process, and not yourself, and forgetting to reset the interested variable at the end of the process.

(d) [Bookwork]

First-fit: allocate first memory section that is big enough. Advantages: fast allocation method.

Disadvantages: can be very inefficient

Best-fit: allocate the smallest section that is enough. Advantages: less inefficient in terms of space than first-fit. Disadvantages: tends to produce lots of remaining tiny fragments, and requires search through the entire list of memory sections

Worst fit: allocate the largest section that is available. Advantages: after allocating the request, the remainder of that section might still be usable. Disadvantages: requires search through the entire list of sections

[1 point per correct algorithm, including its advantages/disadvantages] [3]

After marking comments: most people that attempted this did it correctly, with a minority forgetting some of the advantages or disadvantages, with no particular error pattern.